



# Eptesicus Laboratories

Business Pitch Text Pack

Copy blocks for investor outreach and verbal pitches. Company-first; Lumis-1 as the initial wedge.

## One-liner

Eptesicus Laboratories builds a governance and evaluation runtime that makes on-device AI deployments more auditable and cost-predictable, starting with Lumis-1 as the first packaged product.

## 30-second pitch

Enterprises want AI at the edge for privacy, latency, and cost control, but reliability and governance stop deployments from scaling. Eptesicus Laboratories provides a model-agnostic runtime that wraps small models with validators, calibration, and an evaluation harness. Lumis-1 is the first product instance: an on-device assistant package with a Validator Council and a bounded steering loop. We are raising €50,000 to build reliability datasets, harden the runtime, and ship pilot-ready packaging with trace artifacts.

## 2-minute pitch

On-device AI is moving from experimentation to procurement because it offers privacy, low latency, and predictable costs. In practice, teams hit the same barrier: they can run a model on-device, but they cannot govern it in a way that is auditable, testable, and stable across releases. Reliability issues show up as unsupported claims, inconsistency, and unclear confidence behavior - and without evidence and traceability, deployments stall.

Eptesicus Laboratories builds the missing layer: a governance and evaluation runtime for on-device AI. It orchestrates a draft-and-score loop, applies calibration and policy controls, and produces trace artifacts and regression reports so that teams can validate measurable deltas versus a baseline and ship versioned releases.

Our first commercialization wedge is Lumis-1, a packaged on-device assistant model plus a lightweight Validator Council (Safety, Consistency, Accuracy/Support) and a bounded steering loop. Lumis-1 proves the platform thesis and provides a concrete product for early enterprise pilots. We are raising €50,000 to build reliability datasets and acceptance tests, harden the runtime, and deliver a pilot-ready release that can convert into annual licenses.

## Market snapshot

Public analyst forecasts (including Gartner, IDC, MarketsandMarkets, and Grand View Research) point to rapid growth in edge AI software and increasing spend on AI governance. Aggregated signals include Edge AI software growing from roughly \$11.8B (2024) to about \$56.8B (2029) and AI governance software reaching approximately \$15.8B by 2030. The practical thesis: as adoption moves from pilots to production, governance and evaluation become budget line items because they reduce operational risk and accelerate approvals.

## Why now

- Enterprise adoption has moved from early excitement to scrutiny: pilots must produce measurable returns and controlled risk profiles.
- Privacy, data sovereignty, and procurement requirements favor deployments with local control and audit artifacts.
- Endpoint hardware and on-device tooling have improved, making real on-device deployments viable for more workflows.

## Business model

- Annual enterprise licenses for the runtime and evaluation harness, packaged with model-specific validators and update rights.
- Paid support and SLA tiers for production deployment, calibration, and release discipline.
- OEM/ISV partnerships for bundling into devices and vertical software products.
- Selective paid onboarding and integration, scoped as productized pilots rather than open-ended services.

## Go-to-market

- Start with productized pilots: one workflow, defined acceptance tests, and trace artifacts for evaluation and approvals.
- Convert pilots into annual licenses by expanding workflow coverage and shipping versioned releases.
- Land OEM/ISV bundling where embedded distribution and procurement favor packaged components.
- Position as model-agnostic governance across Apple, Android, and Windows to reduce vendor lock-in.

## Why we win

- Governance and evaluation are shipped as a coherent product: runtime plus validators plus acceptance tests plus trace artifacts.

- Measurable deltas: focus on regression suites and bounded behavior instead of anecdotal demos.
- On-device first: designed for latency budgets, offline constraints, and predictable costs.
- Model-agnostic approach: portable across stacks and avoids single-vendor dependence.

## Risks and mitigations

### Wrapper risk ("just a prompt")

Mitigation: emphasize runtime orchestration, validator models, calibration thresholds, evaluation harness, and trace artifacts as the IP. Package deliverables as versioned software with acceptance criteria and reporting.

### Over-conservatism (over-refusal)

Mitigation: calibration work focuses on reducing unsupported claims while preserving usefulness; thresholds are tuned using failure-mode datasets and measured on regression suites.

### Latency and device budgets

Mitigation: bounded retries, lightweight validators, and profile-driven configuration; ship profiles per device class and workflow.

### Enterprise procurement cycles

Mitigation: productized pilots with clear acceptance tests and trace artifacts; convert to licenses with support or SLA once value is demonstrated.

## Funding ask: €50,000

Raising **€50,000** to build reliability datasets and acceptance tests, harden the runtime and trace artifacts, and deliver a pilot-ready Lumis-1 package.

## Team

- **Deyan Todorov** - Founder
- **Iliyan Bozhanov** - Co-founder

## Status and notice

Eptesicus Laboratories is a brand operated by Deyan Todorov and Iliyan Bozhanov. It is not currently a registered legal entity. The team intends to incorporate prior to signing institutional investment documents. This material is provided for discussion purposes only and is not a product specification or legal advice. It is not intended for medical, legal, or safety-critical decision-making.