

# Lumis-1 Pitch Pack

On-device assistant model with validator-guided steering (offline)

Date: 2025-12-15

Owner: Deyan Todorov (Bulgaria, EU)

Brand: Eptesicus Labs

Contact: [deyan.todorov21@gmail.com](mailto:deyan.todorov21@gmail.com) | +359 89 655 3406

This is a practical, non-hype pitch. It describes what exists, what will be built next, and how success is measured.

## 1) What Lumis-1 is

Lumis-1 is a small on-device assistant model that corrects itself before it answers. It is built from three parts:

- **Lumis-1 Model (Speaker)**: the assistant model you fine-tune for your app.
- **Lumis-1 Validators (Council)**: three lightweight scorers (Safety, Consistency, Accuracy).
- **Lumis-1 Runtime**: an orchestrator that runs draft -> score -> steer -> retry (up to N) -> final answer plus trace.

The point is steering. When a draft is weak, inconsistent, or likely unsupported, the runtime feeds a short correction instruction back into the model and retries.

## 2) Current steering policy

- Unsafe content -> **REFUSE**.
- Otherwise -> always return an answer.
- If Consistency or Accuracy flags issues -> **STEER and RETRY** (max **N = 3** attempts).
- If still not clean -> return the best answer plus a **LOW\_CONFIDENCE** label.
- For external facts offline -> generic answer plus uncertainty (avoid inventing specifics).

### **3) How it works (high level)**

Pipeline:

- Prompt -> Speaker draft (FAST).
- Council scores the draft (Safety, Consistency, Accuracy).
- Runtime decides: pass, refuse, or steer-and-retry.
- Steering prompt is short and strict.
- Stop after 3 retries. Return final answer plus label plus trace.

#### **Validator implementation note:**

In the end state, validators are lightweight discriminative models. In early demos, they can be approximated using short classification prompts on the same base model.

#### **Speaker vs validator coupling:**

For the first release, validators are calibrated to the Lumis-1 Model (Speaker). Interfaces stay generic so portability can be tested later.

## 4) What success looks like (measured)

The system is judged by measurable improvements over the same base model without steering.

- **Repair success rate:** when a draft is flagged, how often does the loop produce a final answer that passes within N attempts?
- **Unsupported-claim rate:** how often does it invent specific details when it cannot verify offline?
- **Refusal precision:** on unsafe prompts, how often does it refuse correctly?
- **Latency per answer:** fast-pass latency vs steering-loop latency (up to 3 attempts).

## 5) Funding ask (EUR 50k)

Use of funds:

- Datasets for the Lumis-1 model and validator council (negatives, contradictions, uncertainty behavior).
- Training and evaluation loops (repeatable runs with reports).
- Development workstation and test devices (non-phone). Exact models are intentionally omitted in first contact.

What you get for this round:

- A stable steering runtime with traces.
- A first model fine-tune aligned with steering behavior.
- An Android application prototype that demonstrates the same policy on-device.
- A reproducible eval report.

## **6) Roadmap (targets)**

Targets with acceptance tests (not calendar promises). Expect +/- 2 weeks variance depending on data generation and training stability.

### **Weeks 1-4**

- Offline demo runner + traces.
- Eval suite (initial) (30 prompts) and report (initial).

### **Weeks 5-8**

- Validator council trained and calibrated.
- Speaker fine-tune aligned with uncertainty discipline and repair behavior.
- Before/after metrics against base model.

### **Weeks 9-12**

- Android application prototype running the same policy on-device.
- Benchmark notes (latency, memory, battery) on non-phone test devices.